

SCORING, SUBSTITUTION MATRICES AND SCORE AGGREGATION FOR TCR ANTIGEN-SPECIFICITY SEARCH

Mikhail Shugay^{1,2,3*}

¹Institute of Translational Medicine, Russian National Medical State University, Moscow, Russia

²Department of Genomics of Adaptive Immunity, Shemyakin–Ovchinnikov Institute of Bioorganic Chemistry
RAS, Moscow, Russia

³Institute of Molecular Biology NAS RA, Erevan, Armenia

vdjmatch technical appendix

Abstract

`vdjmatch` annotates T-cell receptor (TCR) clonotypes by antigen specificity through fuzzy CDR3 search against a labelled database (VDJdb), built on the `seqtree` search core. This note documents the *scoring* layer: how CDR3 similarity is scored, the data-driven TCR substitution matrix (VDJAM) and why it is region-structured, and how per-hit scores are aggregated into a control-calibrated significance. CDR3 loops resist the assumptions behind BLOSUM/PAM — they cannot be multiply aligned across clonotypes of different length, and germline-encoded flanks (CASS...) are so over-conserved that a naive matrix conflates germline bias with selection. We give a position-debiased log-odds estimator that isolates the substitution signal, show empirically that this signal lives in the non-template (NDN) core while the V/J flanks are germline-determined, and demonstrate the consequences on worked examples (GIL, NLV) using the alignments and CIGAR strings produced by `seqtree`. The E-value theory used for significance is derived in the companion `seqtree` appendix; here we summarize it and give the paired-chain extension.

1. THE SEARCH-AND-SCORE FRAMEWORK

Given a query CDR3 q and a labelled reference set D (VDJdb), `vdjmatch` retrieves the neighbours of q within a fixed scope $\theta = (s, i, d, t)$ (max substitutions, insertions, deletions, total edits) using the `seqtree` engine, which defines a ball $B_\theta(q) = \{x : \text{edit}(q, x) \leq \theta\}$ and returns, per hit, the edit decomposition and a matrix-weighted alignment. The matrix enters as a *ranking* score inside the ball; the ball itself (an edit budget) bounds the candidate set. Each hit carries its antigen labels, an extended CIGAR ($=/X/I/D$), and an alignment score. Significance comes from comparing the neighbour count to a matched background control (§3).

*Correspondence: mikhail.shugay@gmail.com

2. SUBSTITUTION SCORING AND THE VDJAM MATRIX

How BLOSUM IS BUILT, AND WHY IT DOES NOT TRANSFER.. BLOSUM is a log-odds matrix estimated from the BLOCKS database of *ungapped multiple-alignment blocks* of conserved protein regions. Its recipe is: (i) take many trusted blocks — columns of homologous, gaplessly aligned residues; (ii) cluster sequences within a block above a percent-identity threshold and down-weight within-cluster pairs so abundant near-duplicates do not dominate; (iii) count, over all blocks, how often residues a and b co-occur in the *same alignment column*, q_{ab} ; (iv) score $s(a, b) = \frac{1}{\lambda} \log(q_{ab}/p_a p_b)$ against the *global* background frequencies p_a . The signal is “which residues substitute for one another within conserved, column-homologous positions,” and the background is position-independent.

CDR3 loops break every step. (i) There is no trusted multiple alignment: same-specificity CDR3s differ in length and have no column homology, so “alignment columns” do not exist. (ii) The germline-encoded ends are over-conserved — essentially every human TRB CDR3 begins CASS and ends in a short J motif — so a *global* background p_a is wrong: an alanine in the N-terminal CASS is germline-fixed and carries no specificity signal, while an alanine in the centre does, yet BLOSUM would score them identically. A single global matrix estimated over all positions thus conflates germline conservation with antigen-driven selection.

WHAT VDJAM CHANGES (POINT BY POINT).. VDJAM keeps the log-odds idea but replaces the two broken ingredients. (a) *Observation unit*: instead of columns of a multiple alignment, the unit is a *same-antigen single-substitution pair* — two CDR3s annotated to the same epitope, equal length, differing at exactly one position; this is the only “aligned residue pair” obtainable without a multiple alignment, and (like a BLOCKS column pair) it records two residues that are interchangeable without changing function (here, specificity). (b) *Background*: instead of a global p_a , a *position-specific* background $\text{bg}(\cdot | L, p)$ (the residue frequency at that length and position) is used, so over-conserved columns are normalized away analytically rather than masked by hand. (c) *Redundancy*: BLOSUM’s percent-identity clustering is replaced — partly by the position-specific background (germline-shared residues are “expected” and contribute little), and prospectively by the control-calibrated significance layer (§3) that already discounts convergent public clones; per-clonotype clustering can be added exactly as in BLOSUM if needed. (d) *Scope*: BLOSUM is one global matrix; VDJAM is estimated per chain and per region, because (next paragraph) only the NDN core carries a transferable substitution signal.

OBSERVATION UNIT AND ESTIMATOR.. We replace “alignment columns” with *same-antigen single-substitution pairs*: ordered pairs of CDR3s annotated to the same epitope, of equal length, differing at exactly one position — the only clean substitution signal extractable without a multiple alignment. For a substitution $a \leftrightarrow b$ observed at position p in a CDR3 of length L on chain c , let $\text{bg}(\cdot | L, p)$ be the residue frequency at position p among all length- L chain- c CDR3s. With $f(a, b)$ the observed count of $\{a, b\}$ events and

$$e(a, b) = \sum_{\text{event slots } (L, p)} n_{L, p} [\text{bg}(a | L, p) \text{bg}(b | L, p) + \text{bg}(b | L, p) \text{bg}(a | L, p)], \quad \text{VDJAM}(a, b) = \log_2 \frac{f(a, b) + \alpha}{e(a, b) + \alpha}, \quad (1)$$

where $n_{L,p}$ is the number of observed events at (L,p) and α a Laplace pseudocount. The key point is that the background is *position-specific*: at an over-conserved column the residue frequency is peaked, so substitutions there are both rare (small f) and “expected” (matched e), and contribute ≈ 0 to the log-odds — the germline bias is removed analytically rather than by hand-masking. The matrix is consumed by `seqtree` via the squared-distance penalty $\text{pen}(a,b) = s_{aa} + s_{bb} - 2s_{ab}$ (`SubstitutionMatrix.from_similarity`); a dominant uniform diagonal makes higher learned similarity rank a substitution better.

REGION STRUCTURE: NDN IS LEARNED, V/J ARE GERMLINE-DETERMINED.. A CDR3 is V-germline prefix · NDN insert · J-germline suffix. Estimating (1) separately by region shows (Fig. 1) that the substitution signal — the agreement with general amino-acid chemistry, measured as correlation with BLOSUM62 — is concentrated in the NDN core (TRB $r = 0.42$, TRA 0.35), whereas the V flank carries almost none (TRB 0.08, TRA 0.07): V-flank “substitutions” are dominated by germline allele differences, not free chemistry. This motivates a two-part model: **VDJAM is learned for the NDN core**, while the **V and J flanks are auto-computed** from the germline V/J sequences and their trimming. Concretely, a CDR3 position derived from V- (or J-) germline is near-invariant when it survives trimming and reverts to the random-insert background when trimmed; so the V/J contribution is a per-position conservation profile $w(p) = \text{Pr}[\text{position } p \text{ is germline-retained}]$, determined by the gene’s trimming-length distribution, rather than a substitution matrix learned from data. Combined, the NDN matrix scores the core and $w(p)$ down-weights substitutions at germline-retained flanks (realized through a per-position weighting of the base matrix). The germline-retention profile $w(p)$ is computed from the OLGA recombination model (per-gene V/J deletion distributions and germline CDR3-portion segment lengths, via `mirpy’s mir.basic.trimming`) and bundled with `vdjmatch`; Fig. 2 shows it drops monotonically from each anchor (the CASS V-stem and the J motif are germline; the core is insert). Region-aware scoring weights each CDR3 column by $1 - \text{Pr}[\text{germline}]$ and applies the NDN matrix — implemented in `vdjmatch.match.regions` and evaluated below.

3. SCORE AGGREGATION AND SIGNIFICANCE

PER-HIT TO PER-QUERY.. A region-aware similarity score sums the NDN matrix score over the core and the (down-weighted) flank contributions; this is available as an optional ranking score and generalizes the legacy cloglog generalized-linear aggregation over V/CDR1/CDR2/J/CDR3.

CONTROL-CALIBRATED E-VALUE (PRIMARY).. Immune repertoires are biologically redundant, so an i.i.d. null over-calls. `vdjmatch` instead counts a query’s VDJdb neighbours within the ball and compares to the count expected from a matched background control: with $N = |D|$, M the control size and n_C the control neighbour count, $\mathbb{E} = (N/M) n_C$ is the expected number of VDJdb neighbours by chance and $p_{\text{enrich}} = \mathbb{P}(\text{Poisson}(\mathbb{E}) \geq n_D)$ is the enrichment significance — significant only when a clonotype has *more* VDJdb neighbours than the generative background predicts. The Poisson/Chen–Stein theory, the punctured (self-excluding) null and the choice of M are derived in the companion `seqtree` appendix.

PAIRED $\alpha\beta$ E-VALUE.. The null factorization is not an approximation of convenience: α and β chains pair *almost without constraint* in the repertoire at large, as shown by integrating structural




Figure 1: Per-region agreement of the learned CDR3 substitution matrix with BLOSUM62 (Pearson r of off-diagonal scores), human VDJdb. The free-substitution signal lives in the NDN core; the V flank is germline-determined.

and paired-sequencing data [4], so under the background (generative) law the joint ball mass genuinely factorizes, $\pi_0^{\alpha\beta} \approx \pi_0^\alpha \pi_0^\beta$. Crucially, the *same* analysis finds that *antigen-driven selection distorts* this uniform pairing landscape — i.e. conditioned on an epitope the two chains are *not* independent. This is exactly the structure the joint E-value exploits: the null factorizes (background pairing is unconstrained), so among the N paired VDJdb complexes the expected number of chance joint matches is $\mathbb{E}_{\alpha\beta} = N \pi_0^\alpha \pi_0^\beta$ with p the Poisson tail on the count of complexes matched in *both* chains; an excess over this factorized null is precisely the epitope-conditional dependence (co-selection of α and β) that signals a true specificity. Because the joint null is far smaller than either chain’s, true pairs are dramatically more significant: on the paired TCRvdb benchmark, 296/300 pairs match a complex and 295/296 (99.7%) recover the correct epitope, with joint p down to 3.6×10^{-9} . (When the two chains’ hits are treated as independent tests, the equivalent statistic is Fisher’s combination of the per-chain enrichments; the appendix `seqtree` co-occupancy term quantifies the residual dependence.)

4. RESULTS: DOES VDjam (AND REGION WEIGHTING) HELP?

A matrix derived from same-antigen pairs must be tested on epitopes it never saw, or the evaluation is circular. For each held-out epitope E^* we re-derive the NDN matrix from *all other* epitopes, fix the candidate set with a single unit-cost search on E^* (exact self excluded), and re-score each (query, hit) pair four ways — unit, BLOSUM62, NDN-VDjam, and region-weighted NDN-VDjam — ranking by score and reporting micro PR-AUC for “hit shares E^* ” (Fig. 3, eight largest TRB epitopes). Two honest findings. (i) *Region weighting is a real, consistent, but small improvement over the flat matrix* (+0.004 at scope 2, +0.007 at scope 3; never worse on any epitope, and the gain grows with scope) — modest because CDR3 variation is *already* NDN-concentrated, so down-weighting the rarely-varying germline flanks moves few hits; it should matter more for cross-V-gene comparisons, where flank (germline-allele) differences are common. (ii) *BLOSUM62 remains the strongest substitution*




Figure 2: OLGA-derived germline-retention $w(p)$ (via *mirpy*): probability that each CDR3 offset from the V-Cys (or J) anchor is germline-encoded. It falls monotonically into the insert core, defining the region weight $1 - w(p)$ used in scoring.

matrix at this narrow-scope retrieval task (mean PR-AUC 0.32 vs 0.28 region-VDJAM vs 0.29 unit): the data-derived NDN matrix captures general substitution tolerance that BLOSUM already encodes well, and with limited same-antigen pairs it is noisier. At scope 2 the search ball supplies most of the specific-vs-nonspecific discrimination and any matrix only re-ranks within it; the substitution model is therefore a second-order lever here, and the first-order signal is the control-calibrated neighbour count (§3).




Figure 3: Leave-one-out-by-epitope retrieval PR-AUC (TRB, scope 2): the NDN VDJAM matrix versus BLOSUM62 and unit cost on epitopes excluded from matrix estimation.

5. EDIT DISTANCE, POSITION, AND THE MATRIX COMPARISON

Four analyses (`bench/scoring_analysis.py`, `bench/loo_vdjam.py`) settle what scoring actually buys on VDJdb. All figures and numbers below are the *human TRB* example case (the largest, best-populated slice of VDJdb); the same pipeline applies unchanged to human/mouse and to TRA/TRB — point `$VDJDB_SAMPLE/$VDJDB_SPECIES` at another export to regenerate.

HAMMING DISTANCE 1 IS THE SIGNAL:NOISE OPTIMUM.. The fraction of CDR3 neighbours that share an epitope — the precision of a similarity call — falls steeply with edit distance: macro purity (per-epitope mean, robust to the $100\times$ epitope-size range; § below) 0.44, 0.23, 0.14, 0.10, 0.07 at distance 1 ... 5, an enrichment of $44\times \rightarrow 2.3\times$ over chance (Fig. 5). A distance-1 neighbour is twice as likely to be a true co-specific as a distance-2 one and six times a distance-5 one. This reproduces the original VDJdb observation [6] that single-mismatch neighbours are the sweet spot — and on the dense 2026 release it survives only once a few 10x mega-studies are capped and the random tail is treated as an admixed control (§3); the naive pooled rate reads a near-constant ~ 0.9 because a random clonotype is probably from a 30k-clonotype epitope. It frames everything below: the search ball, not the substitution matrix, is the dominant lever; the matrix only re-ranks within a ball whose purity is already set by its radius.

CENTRAL SUBSTITUTIONS CARRY THE SPECIFICITY SIGNAL.. Conditioning on *where* a single substitution falls (Fig. 6), the probability that the two CDR3s still share an epitope dips to ≈ 0.31 in the NDN core but rises to ≈ 0.90 toward both the V and J borders: a central mismatch most often *changes* specificity (it is contact-proximal and informative), while a border mismatch is usually germline noise (the conserved V/J stems). This is the empirical basis for weighting the centre more — the same direction as the germline-retention weight $1 - w(p)$ (§2) — and, encoded directly as a positional weight on BLOSUM62 (the PSSM of §5), it is the one change that clearly improves retrieval over flat BLOSUM62 (7/8 epitopes; Table 1). It also matches the contacting-region picture of Shcherbinin et al. [5].

THE NDN CORE IS GLYCINE-RICH.. Central CDR3 positions reach $\sim 31\text{--}34\%$ glycine in both chains (Fig. 7) — the N-region/D-gene insertion signature — a strongly non-BLOSUM composition. Notably TRA is comparably central-G-rich to TRB, so this is an insertion-diversity feature, not a TRB-D-gene-exclusive one; either way the NDN alphabet is distinctive enough that a general matrix is a poor fit there in principle.

NO AMINO-ACID MATRIX CLEARLY BEATS BLOSUM62 — AND A GENETIC-CODE NULL TIES IT.. Re-scoring the held-out-epitope candidates (leave-one-out, § above) with each substitution matrix, scored by balanced PR-AUC on the composition-controlled long list, gives a consistent verdict at every scope (Table 1, Fig. 4): BLOSUM62, PAM250 and the structural matrix are essentially tied at the top, plain Hamming (unit) is close behind, and the data-derived NDN-VDJAM — flat or region-weighted — trails them all. Strikingly, **VDJAMr**, a matrix built from the *genetic code alone* (`loo_vdjam.codon_dissim`: $M(a, b) = \Pr[\text{a single random nucleotide substitution of a random codon of } a \text{ yields } b]$, with no chemistry or selection), *matches* BLOSUM62 (+0.013 at distance ≤ 2 , +0.016 at ≤ 4 ; ahead in 4–5/8 epitopes) and clearly beats the data-derived VDJAM. That a pure codon-mutational-accessibility null does as well as a chemically tuned matrix says the substitution structure of co-specific CDR3s

is largely *generative* (what the recombination/mutation process reaches cheaply), not chemical — which is exactly why hand-tuning the amino-acid alphabet does not help. The substitution matrix is a second-order lever; the first-order statistic remains the control-calibrated neighbour count at distance 1 (§3).

BUT POSITION-WEIGHTING BLOSUM62 DOES (THE SIGNIFICANCE PSSM).. Experiment 2 says *where* a substitution falls matters more than *which* substitution it is. We encode that directly as a positional factor on top of BLOSUM62 rather than as a new alphabet matrix. From the position-significance curve (Fig. 6) we take a per-relative-position weight $\omega(p) = (1 - \text{Pr}(\text{same} \mid \text{sub at } p))$, normalised to mean 1 (centre ≈ 1.9 , V/J borders $\approx 0.2\text{--}0.4$), and interpolate it to any CDR3 length. `seqtree` consumes this natively: `regions.significance_pssm()` builds a fixed-width `PositionalMatrix.from_weights(BLOSUM62 [100  $\omega(p)$ ])`, so the engine scores $\text{pen}(p, a, b) = \omega(p) \cdot \text{BLOSUM62}(a, b)$ — a central mismatch costs $\sim 2\times$ its border-position cost. This is the only scoring change that clearly beats flat BLOSUM62: $0.572 \rightarrow 0.623$ at distance ≤ 2 (winning 7/8 epitopes) and $0.589 \rightarrow 0.617$ at ≤ 4 (6/8) (Table 1). The gain is uniformly signed because the matrix re-ranks within a ball whose purity is already set by the radius. The remaining lever, the *V-gene* dimension, is taken up in §6.

A PURPOSE-BUILT TCR MATRIX (TCRBLOSUM) DOES NOT TRANSFER.. Postovskaya et al. [3] report that `tcrBLOSUM` — a BLOSUM-style matrix counted from same-epitope, equal-length CDR3 pairs — improves epitope-specific TCR clustering over BLOSUM62. Two features make that gain suspect: the counts apply *no* redundancy clustering (so convergent public clonotypes dominate the statistics), and the matrix is validated on the same same-epitope relation it is fit to. Run through our leave-one-out-by-epitope retrieval (their published `tcrBLOSUMb`, `bench/tcrblosum_refute.py`), it does *not* beat BLOSUM62: 0.557 vs 0.570 at distance ≤ 2 (ahead in 3/8 epitopes), 0.580 vs 0.586 at ≤ 4 (5/8), while their bundled BLOSUM62 reproduces ours exactly (a loader sanity check). Reweighting *either* matrix by the central-significance PSSM helps, and it helps BLOSUM62 the most (BLOSUM+possig 0.622 vs `tcrBLOSUM`+possig 0.617). The lesson is consistent: a re-counted amino-acid alphabet does not transfer across epitopes; position does.

balanced PR-AUC	unit	BLOSUM62	PAM250	structural	VDJAM	VDJAMr	VDJAM-region	BLOSUM+possig
edit distance ≤ 2	0.545	0.572	0.575	0.572	0.519	0.585	0.525	0.623
edit distance ≤ 4	0.576	0.589	0.597	0.589	0.535	0.605	0.544	0.617

Table 1: Substitution matrices by leave-one-out-by-epitope VDJdb-vs-VDJdb retrieval (human TRB, 8 largest epitopes of the 2026-06-11-ZENODO composition-controlled long list; balanced PR-AUC). The chemical matrices (BLOSUM62/PAM250/structural) and the genetic-code null VDJAMr are all $\approx$ tied; the data-derived VDJAM trails; only BLOSUM62 reweighted by the positional-significance PSSM clearly leads (7/8 and 6/8 epitopes).

6. THE V-GENE DIMENSION

The remaining lever is the V gene. Three questions on the composition-controlled long list (human TRB unless noted): is a V match itself a specificity signal, does it soften to germline-similar V genes, and — if so — does it localise to a few contacting-loop positions (a V “pseudosequence”)?




Figure 4: Mean leave-one-out retrieval PR-AUC by scoring (TRB, distance ≤ 2). No standard amino-acid matrix beats BLOSUM62; reweighting BLOSUM62 by the central-position significance factor (green) does.

A V MATCH IS A STRONG, NEAR-BINARY SPECIFICITY PRIOR.. Splitting every scored neighbour pair by whether query and hit use the same V family, same-V neighbours share the held-out epitope 44% of the time versus 6% for cross-V (edit distance ≤ 4 ; Table 2) — a $\sim 7\times$ prior, far larger than any substitution-matrix effect. This is why the legacy tool used V as a hard filter, and it argues for carrying V as a strong prior (a soft $S(V)$ term or a per-V null) in a joint V+CDR3 score. Position weighting sharpens the *same-V* stratum most (balanced PR $0.632 \rightarrow 0.713$ with the significance PSSM, $+0.081$) — where the germline frame is shared and only the NDN core varies — and little in the noise-dominated cross-V stratum ($+0.012$).

COARSE GERMLINE-LOOP SIMILARITY BARELY PREDICTS IT.. If the same-V advantage were simply the chemistry of the V-encoded contacts, cross-V co-specificity should rise smoothly with germline CDR1+CDR2 similarity. As a *continuous* quantity it barely does: across every $\{\text{human, mouse}\} \times \{\text{TRA, TRB}\} \times \{\text{MH I, MHC-II}\}$ cell (`bench/vgene_scan.py`, Fig. 8, Table 4) the V prior is universal (same/cross ratio $1.0\text{--}8\times$) but the point-biserial correlation between loop similarity and co-specificity is ≈ 0 ($r \in [-0.06, +0.06]$ in the well-powered cells). Binning cross-V pairs by similarity shows only a weak rise — 6% at low similarity to 12% at the top bin, still far below the same-V 44% (Table 3). So a loose, continuous “fuzzy-V” similarity does not recover the prior.

BUT NEAR-EXACT GERMLINE IDENTITY DOES — AT CDR3-LIKE TOLERANCE.. Resolving the similarity to exact germline-region matches reverses that impression (`bench/vregion_decompose.py`, human TRB MHC-I; Table 5). Cross-V co-specificity rises *monotonically* as the germline CDR1+CDR2 edit distance shrinks — 10.9% at ≥ 6 mismatches, 13.7% at 3–5, 20.6% at 1–2, and 60.0% at 0 (identical loops), approaching the same-V level (65%). FR1 identity behaves likewise (50%, $4.4\times$). So the V signal *is* carried by the germline contacting loops, but with a tight, CDR3-like tolerance: it needs near-exact identity, which is why the continuous metric above (dominated by 1–2-mismatch pairs) misses it. Caveat: the exact-match bins are small ($n = 25$ at distance ≤ 2 , $n = 4$ at ≤ 1); the claim rests on the monotone trend across the well-powered bins ($n = 568, 3679, 34095$), with




Figure 5: Macro purity (per-epitope mean fraction of CDR3 neighbours sharing an epitope) vs Hamming distance (human TRB, composition-controlled long list): distance 1 is the signal:noise optimum [6].




Figure 6: $\Pr(\text{neighbour shares epitope} \mid \text{single substitution at relative position})$, human TRB. The NDN core is where a mismatch most changes specificity; the V/J borders are germline noise.

the identical-loop endpoint suggestive.

THE RECOVERY IS WHOLE-LOOP, NOT A SPARSE PSEUDOSEQUENCE.. Do a few decisive positions carry it (the MHC-pseudosequence analogue)? Per-position residue-match lift over the cross-V baseline (`bench/vpseudo.py`, length-constant CDR1/CDR2) is mostly flat (lift 1.0–1.2 at the 10.5% baseline), with a single exception — CDR2 position 5 ($1.67\times$, $n = 5392$). So the prior is recovered by *whole-loop* near-identity, not by matching a handful of pseudosequence coordinates: a different V gene must reproduce essentially the entire CDR1+CDR2 to behave like the same V. A useful soft-V match must therefore demand near-exact loop identity, not loose similarity. The next step is the full IMGT-numbered treatment (FR positions and the FR3 DE-loop/“CDR2.5”, via the `arda` IMGT/V-QUEST milestone) to test whether specific positions sharpen this.




Figure 7: Glycine fraction by relative CDR3 position (human): the NDN core is $\sim 30\text{--}36\%$ glycine in both TRA and TRB (insertion/D-gene signature).

V stratum	pairs	co-specific	flat	+possig	+region
same V	116 919	44.1%	0.632	0.713	0.680
cross V	981 672	6.3%	0.560	0.573	0.537

Table 2: V-gene-stratified leave-one-out retrieval (human TRB, edit distance ≤ 4 ; co-specific = positive rate, the rest balanced PR-AUC). A V match is a $\sim 7\times$ co-specificity prior; the significance PSSM helps the same-V stratum most (+0.081).

7. BENCHMARK: THE REFERENCE-REPLICATED SHORTLIST

The latest VDJdb release is our standing benchmark, but its labels vary in confidence. The cleanest gold standard is a clonotype–epitope association reported independently in ≥ 2 references (`db.replicated`, null references excluded; ~ 3000 TRB and ~ 2000 TRA such pairs in the current human release) — the same “shortlist” idea used for `mhcmatch`. We measure annotation accuracy on it by leave-one-out nearest-neighbour voting against the rest of VDJdb at the signal:noise-optimal scope (subs 1), excluding the clonotype’s own exact CDR3 so the label must come from *other* TCRs (`bench/shortlist_accuracy.py`). Replicated TCRs are annotated at $\sim 44\text{--}47\%$ top-1, modestly above size-matched single-reference controls of the same epitopes (Table 6). The dramatic $3\times$ gap we saw on a sparse older export was an artefact: on the dense 2026 release even single-reference clonotypes have findable neighbours (control recall $\sim 48\text{--}72\%$), so the gold set is cleaner but no longer dramatically *easier*. Restricting neighbours to the same V family helps at the wider subs 2 scope (TRB $53.1 \rightarrow 56.8\%$) — exactly where cross-V false neighbours appear (§6), confirming the V prior is worth its place.

8. WORKED EXAMPLES

PAIRWISE (DIFFERENT LENGTHS).. `seqtree` aligns CDR3s of differing length, placing indels where the loop genuinely lengthens. The edit-distance alignments below (influenza GIL-specific TRB CDR3s) show the conserved CASS V-anchor and DTQYF J-motif with substitutions and indels confined to the central NDN core:

germline CDR1+CDR2 similarity	cross-V pairs	co-specific
[0.0, 0.3)	484 238	6.0%
[0.3, 0.5)	385 956	6.3%
[0.5, 0.7)	87 913	7.0%
[0.7, 0.9)	19 533	9.7%
[0.9, 1.0)	4 032	12.3%

Table 3: Continuous germline CDR1+CDR2 similarity gives only a weak rise in cross-V co-specificity (6% $\rightarrow$ 12%), far below the same-V 44% — which is why $r \approx 0$. The strong signal lives at near-*exact* identity (Table 5), not along the continuous similarity axis.




Figure 8: Same-V vs cross-V CDR3-neighbour co-specificity across species/chain/MHC-class cells of VDJdb (h/m = human/mouse; A/B = TRA/TRB; I/II = MHC class). A V match is a universal co-specificity prior; continuous CDR1/CDR2 similarity does not bridge the gap (Table 4).

query (ref): CASSPRSGETQYF epitope GILGFVFTL (unit / edit-distance alignment)

vs CASSNRSGETQYF (len 13; subs=1 ins=0 dels=0; CIGAR 4=1X8=)

```
CASSPRSGETQYF
||||| |||||
CASSNRSGETQYF
```

vs CASSIRSGESTQYF (len 14; subs=1 ins=1 dels=0; CIGAR 4=1X4=1I4=)

```
CASSPRSGE-TQYF
||||| |||||
CASSIRSGESTQYF
```

vs CASSLFSETQYF (len 12; subs=2 ins=0 dels=1; CIGAR 4=2X1=1D5=)

```
CASSPRSGETQYF
||||| | |||||
CASSLFS-ETQYF
```

cell (species / chain / class)	same-V	cross-V	ratio	$r(\text{vsim, co-spec})$
human TRA MHC-I	57.9%	17.3%	3.3	0.060
human TRA MHC-II	98.0%	24.8%	4.0	0.242
human TRB MHC-I	64.3%	8.0%	8.0	0.056
human TRB MHC-II	97.5%	76.1%	1.3	0.046
mouse TRA MHC-I	75.1%	29.3%	2.6	0.021
mouse TRA MHC-II	50.4%	48.0%	1.0	-0.010
mouse TRB MHC-I	76.7%	40.6%	1.9	-0.030
mouse TRB MHC-II	83.7%	31.6%	2.7	-0.016

Table 4: Per-cell same-V vs cross-V co-specificity and the point-biserial correlation between continuous germline CDR1+CDR2 similarity and co-specificity among cross-V pairs (top 6 epitopes/cell, ≥ 30 CDR3/epitope). The V prior holds everywhere (ratio 1–8 $\times$); continuous loop similarity barely predicts ($r \approx 0$; the 0.24 is human TRA MHC-II, whose high-similarity bin has $n = 4$).

cross-V stratum (human TRB MHC-I)	co-specific	n
<i>cross-V baseline</i>	11.4%	38 419
CDR1+CDR2 edit distance ≥ 6	10.9%	34 095
edit 3–5	13.7%	3 679
edit 1–2	20.6%	568
edit 0 (identical)	60.0%	25
FR1 identical	50.3%	338
<i>same-V reference</i>	65.2%	12 869

Table 5: Cross-V co-specificity rises monotonically as germline CDR1+CDR2 approach identity, reaching $\approx$ the same-V level at exact identity (edit 0); FR1 identity behaves likewise. Near- *exact* germline-loop identity — not loose similarity — recovers the V prior (edit 0/CDR1&CDR2 bins small, $n = 25$).

vs CASSLGETQYF (len 11; subs=1 ins=0 dels=2; CIGAR 4=2D1X6=)

CASSPRSGETQYF

|||| |

CASS--LGETQYF

vs CASSPWSGSQETQYF (len 15; subs=1 ins=2 dels=0; CIGAR 5=1X2=2I5=)

CASSPRSG--ETQYF

|||| |

CASSPWSGSQETQYF

MULTIPLE (SAME LENGTH).. Aligning same-length same-epitope CDR3s (no indels needed) makes the conservation profile explicit (Figs. 9–10): columns are shaded by conservation, dark at the germline-determined CASS prefix and J motif, light through the variable NDN core. This is the structural counterpart of the finding of Shcherbinin et al. [5] that specificity-preserving CDR3 substitutions occur predominantly outside the peptide/MHC-contacting positions while preserving the loop backbone — i.e. the tolerated variation we score in the NDN core is exactly the variation that

chain	set	n	top-1 (CDR3)	top-1 (V-restr.)	recall
TRA	shortlist (≥ 2 ref)	1993	43.7%	43.7%	62.3%
TRA	control (1 ref)	1993	32.2%	27.9%	48.4%
TRB	shortlist (≥ 2 ref)	2962	47.2%	46.9%	54.2%
TRB	control (1 ref)	2962	45.1%	41.5%	48.8%

Table 6: Leave-one-out epitope-annotation accuracy on the reference-replicated shortlist vs a size-matched single-reference control (human, subs 1, exact self excluded). Top-1 = nearest-neighbour-vote epitope correct; recall = true epitope reachable within scope. On the dense 2026 release the shortlist is modestly above control, not the $3\times$ of the older sparse export.

does not break recognition.

9. LIMITATIONS AND ROADMAP

The substitution *alphabet* is a second-order lever: the data-derived NDN-VDJAM trails BLOSUM62 at every scope, and a chemistry-free genetic-code null (VDJAMr) does just as well as BLOSUM62 — the co-specific-CDR3 substitution structure is largely generative, not chemical (§5). What moves the needle is *position*: reweighting BLOSUM62 by the central-significance PSSM clearly beats flat BLOSUM62 (7/8 epitopes), confirming that for CDR3 *where* a mismatch falls matters more than *which* residue it is. The *V-gene* dimension (§6) is a strong near-binary prior ($\sim 7\times$); it is *not* captured by loose germline-loop similarity, but cross-V co-specificity does rise to $\approx$ the same-V level at *near-exact* CDR1+CDR2 identity — a tight, CDR3-like tolerance, recovered whole-loop rather than via a sparse pseudosequence (exact-match bins small, $n \approx 25$). The open work is to model V as a strong prior in a joint V+CDR3 E-value, to demand near-exact (not loose) loop identity for any soft-V match, and to extend the position analysis to full IMGT numbering (FR positions and the FR3 DE-loop, via the *arda* IMGT/V-QUEST milestone); a per-chain treatment (TRA vs TRB) is also pending. Separately, the control’s V–J usage must match the query repertoire’s or the null is mis-specified — whether to renormalize V–J usage (at the risk of erasing genuine V-gene bias) is an open question to settle on held-out epitopes.

A further direction is to predict, *a priori*, which epitopes are intrinsically easy versus hard to annotate — the wide spread of per-epitope PR-AUC in our experiments (e.g. the convergent, “featured” CMV NLV vs the diffuse, “featureless” influenza GIL) is not noise but biology. A pMHC with prominent TCR-facing features selects a focused, motif-rich repertoire that clusters cleanly, whereas a featureless one is recognised by structurally diverse TCRs in varied docking modes and yields no tight motif [8, 9, 7, 2]; clustering-based specificity inference succeeds or fails accordingly [1]. A model that scores epitope “annotability” from repertoire convergence (and, where available, pMHC structural features) would let us report calibrated confidence per epitope and triage GIL-like hard cases — a planned addition to the benchmark. These extensions, and the full comparison against existing tools, are ongoing work.

GILGFVFTL (TRB CDR3, length 13, n=14)




Figure 9: GIL (influenza M1, A*02) TRB CDR3 block; conserved V/J flanks, variable NDN core.

REFERENCES

- [1] Dan Hudson, Alex Lubbock, Mark Basham, and Hashem Koohy. A comparison of clustering models for inference of T cell receptor antigen specificity. *Immunoinformatics*, 13:100033, 2024.
- [2] Hung-Pin Peng, Kuo Hao Lee, Jhih-Wei Jian, and An-Suei Yang. Origins of specificity and affinity in antibody-protein interactions. *Proceedings of the National Academy of Sciences of the United States of America*, 111(26):E2656–E2665, 2014.
- [3] Anna Postovskaya, Koen Vercauteren, Pieter Meysman, and Kris Laukens. terBLOSUM: an amino acid substitution matrix for sensitive alignment of distant epitope-specific TCRs. *Briefings in Bioinformatics*, 26(1):bbae602, 2024.
- [4] Dmitrii S. Shcherbinin, Vlad A. Belousov, and Mikhail Shugay. Comprehensive analysis of structural and sequencing data reveals almost unconstrained chain pairing in TCR $\alpha\beta$ complex. *PLoS Computational Biology*, 16(3):e1007714, 2020.

NLVPMVATV (TRB CDR3, length 15, n=14)




Figure 10: NLV (CMV pp65, A*02) TRB CDR3 block; same conservation pattern.

- [5] Dmitrii S. Shcherbinin, Vadim K. Karnaukhov, Ivan V. Zvyagin, Dmitriy M. Chudakov, and Mikhail Shugay. Large-scale template-based structural modeling of T-cell receptors with known antigen specificity reveals complementarity features. *Frontiers in Immunology*, 14:1224969, 2023.
- [6] Mikhail Shugay, Dmitriy V. Bagaev, Ivan V. Zvyagin, Renske M. Vroomans, Jeremy Chase Crawford, Garry Dolton, et al. VDJdb: a curated database of T-cell receptor sequences with known antigen specificity. *Nucleic Acids Research*, 46(D1):D419–D427, 2018.
- [7] InYoung Song, Anna Gil, Rabinarayan Mishra, Dario Gherzi, Liisa K. Selin, and Lawrence J. Stern. Broad TCR repertoire and diverse structural solutions for recognition of an immunodominant CD8+ T cell epitope. *Nature Structural & Molecular Biology*, 24(4):395–406, 2017.
- [8] Stephen J. Turner, Katherine Kedzierska, Helen Komodromou, Nicole L. La Gruta, Michelle A. Dunstone, Andrew I. Webb, Richard Webby, Helen Walden, Weidong Xie, James McCluskey, Anthony W. Purcell, Jamie Rossjohn, and Peter C. Doherty. Lack of prominent peptide-major histocompatibility complex features limits repertoire diversity in virus-specific CD8+ T cell populations. *Nature Immunology*, 6(4):382–389, 2005.
- [9] Xinbo Yang, Guobing Chen, Nan-Ping Weng, and Roy A. Mariuzza. Structural basis for clonal diversity of the human T-cell response to a dominant influenza virus epitope. *The Journal of Biological Chemistry*, 292(45):18618–18627, 2017.